

Useful Metacognition for Math Reasoning — Why It Hasn't Worked Yet, and the Fix We Are Testing

Author: Seungpil Lee (+ Claude) **Date:** 2026-06-22 **Project:** metacognition-math (confidence-rv)

Executive Summary

We are training a language model to decide when to think about its own thinking (verify / redirect / abstain) so that doing so **raises** its math accuracy. We taught the model the format successfully, but reinforcement learning (RL) made the metacognition **net-harmful**. This report explains, hypothesis by hypothesis, what we tried, what came out, and what each result means — and the fix we are now implementing.

Hypothesis	What we tested	Headline metric	Result	Verdict
H1	SFT-taught verify/redirect + RL (reward = "PMI") makes meta causally useful	$\Delta = \text{acc_with} - \text{acc_without}$	gs100 -0.013 → gs190 -0.008	Falsified (meta net-harmful, flat with training)
H2	The PMI reward rewards coherence, not causal usefulness →	meta content + emission rate	94% generic "verify", emission 1.000, breaks 16 > saves 11	Confirmed (mechanism)

Hypothesis	What we tested	Headline metric	Result	Verdict
	generic-verify collapse			
H3	Rewarding the counterfactual answer-change fixes it (abstain + don't break correct answers)	—	implementation in progress	Pending

One-sentence conclusion: the model's metacognition fails not because it is badly trained or badly formatted, but because the **reward measures the wrong thing** — it pays the model for writing a plausible-looking "let me double-check," regardless of whether the check ever catches anything; we are replacing it with a reward that pays only when the metacognition actually changes a wrong answer into a right one.

1. Background and Intent (for a first-time reader)

The goal. Humans solving a hard problem do metacognition: "I'm not sure about this step — let me verify it," or "this approach is going nowhere — let me switch methods." We want a model to do the same — but **only when it helps**. Concretely, inside its answer the model emits a special block:

```
<|meta|> ... reasoning about its own confidence ... decision: verify | redirect | none
```

- **verify** — the model is fairly confident and the claim is checkable, so it does an independent check.
- **redirect** — the model thinks it is likely wrong / stuck, so it switches method.
- **none** — no metacognition needed (abstain).

The north-star. Useful metacognition → higher accuracy, and only when useful. Metacognition is a **means**, not an end. Getting the model to talk about its confidence is easy; getting that talk to **causally raise accuracy** is the hard part, and is the entire point.

How we measure "useful" (the key metric). For every test problem we decode the answer **twice** under identical weights:

- **arm A (with meta):** normal decoding — the model may emit `<|meta|>` .
- **arm B (without meta):** the `<|meta|>` open token is forbidden (its logit driven to $-\infty$), so the model must solve without metacognition.

Then $\Delta = \text{accuracy}(\text{with}) - \text{accuracy}(\text{without})$ is the **causal contribution of metacognition**. $\Delta > 0$ means meta helps; $\Delta < 0$ means meta hurts. This is the number every hypothesis below is judged on.

2. Method (the pipeline)

Stage	What it does
Stage-0	A teacher model (GPT-class via TRAPI) distills a corpus of verify/redirect examples anchored on the student's own wrong attempts; the gold answer is used only to check correctness, never to set confidence (no-leak). 562 redirect + 1219 verify rows.
Stage-1	Warm-up supervised fine-tuning (SFT) on that corpus, so the model learns the meta format.
Stage-2	RL (a GRPO-variant we call DCPO) that is supposed to make the meta causally useful, using several reward "heads" combined into the policy-gradient advantage.

The Stage-2 reward combines five heads onto the advantage: **correctness** (is the final answer right), **R_meta** (the meta-quality signal — this is the one that matters here), **calibration**, **emission** (don't go silent), and **format** (well-formed meta). The current R_meta is **PMI** (defined in §4).

3. Results, Hypothesis by Hypothesis (with interpretation)

3.1 H1 — "SFT + RL with the PMI reward makes meta causally useful ($\Delta > 0$)"

Intermediate result (Stage-1, format only). The warm-up SFT passed its gate: accuracy-with-meta rose **+8.6 pp** over the initialization on the same validation set, and the model emitted well-formed meta on essentially every problem. Interpretation: the model **learned the format** and did not lose accuracy — but this says nothing about causal usefulness; it only says "the meta is there and looks right."

Main result (Stage-2, the real test). After RL, the counterfactual Δ was **negative and stayed negative** as training continued:

Checkpoint	acc_with	acc_without	Δ	saved	broke	emission
gs100	0.710	0.724	−0.013	12	20	1.000
gs190 (+90 steps)	0.741	0.749	−0.008	11	16	1.000

(source: HF [iamseungpil/metacot-rv](#) eval/rv_dcpo_gs100, eval/rv_dcpo_gs190; n=594.)

Interpretation — H1 is falsified, two ways: 1. **Meta is net-harmful.** $\Delta < 0$ means turning metacognition off would raise accuracy. The "saved vs broke" columns show why: the meta rescues a handful of wrong answers (11–12) but **breaks more correct ones** (16–20). Net negative. 2. **It is not a "needs more training" problem.** Over 90 extra steps Δ moved only $-0.013 \rightarrow -0.008$ (within noise), while absolute accuracy rose ($0.710 \rightarrow 0.741$ acc_with) — i.e., the model got better at solving but the meta contribution stayed harmful. Because the mechanism (below) does not change with more steps, extrapolating Δ to positive is unjustified. **The bottleneck is structural, not the amount of training.**

Where exactly does meta hurt? (stratified, gs190)

Stratum	n	Δ
easy	198	−0.005

Stratum	n	Δ
medium	261	-0.004
hard	135	-0.022
redirect (scenario)	309	-0.013
verify (scenario)	285	-0.004

Interpretation: there is **no regime where meta is clearly positive** — it ranges from roughly neutral (easy/medium, where it occasionally breaks a correct answer) to clearly harmful (hard, where redirect flails). A useful metacognition system would show at least one stratum with $\Delta \gg 0$; this one does not.

A subtle but important reading (Stage-1 vs Stage-2 dissociation): SFT fixed the format (C-gate +8.6 pp) but RL could not make it causally useful. That dissociation points the finger at the **RL reward**, not at the SFT data or the format — which sets up H2.

3.2 H2 — "The PMI reward rewards coherence, not causal usefulness"

Result (reading what the model actually writes). We read the emitted meta text and counted decisions:

- **~94% of metas are decision: verify** (14 verify : 1 redirect in the sampled stream).
- **emission rate = 1.000** — the model always emits meta; it never abstains.
- The verify is a near-identical **template**: "X looks plausible but should be independently checked before accepting it → decision: verify," followed by a **full re-derivation** of the answer (not a targeted check of one risky step).

Interpretation — confirmed, with the mechanism. On an easy problem the answer is already correct, so a fluent "let me verify by re-deriving" is (a) **coherent** and (b) **lands on a correct answer** — and the PMI reward (§4) pays for exactly that combination. There is **no term anywhere that asks "would the answer have been correct without this meta?"** So unnecessary meta is never penalized, only coherent meta is rewarded → the reward's optimum is "emit a coherent generic verify everywhere," and re-deriving an already-correct answer is a pure downside dice-roll that sometimes recomputes wrongly and **overrides a correct answer** (the "broke 16 > saved 11"). This mechanistically explains H1's flat-negative Δ . (Caveat: this is a triangulation of content + reward

definition + the Δ pattern, not a controlled ablation — which is exactly what H3 provides.)

3.3 H3 — "Reward the counterfactual answer-change, not coherence" (being tested)

The idea. Replace the coherence reward with the thing we actually care about: did the meta **change a wrong answer into a right one**? That is the counterfactual `correct_with - correct_without` — literally the north-star metric used as the reward. If the meta does not change the answer (easy, already correct), the reward is zero → the model is **not paid to emit unnecessary meta** → it learns to **abstain**. If the meta breaks a correct answer, the reward is negative → it learns **not to break**.

Doing it cheaply (the trick from prior work). Naively this needs a second decode per training rollout ($\sim 2\times$ cost), which is why PMI was used instead. The fix (GRPO-as-PRM, 2509.21154) is to **branch the existing rollout group**: we already sample 8 rollouts per problem; we make 4 of them emit meta and force 4 to solve without meta (the same meta-open ban used at eval). The group-relative advantage then equals the counterfactual answer-delta **for free** — no extra decode. (See §5.)

Status: implementation in progress (§6). **Prediction (how we will judge it):** Δ rises above gs190's -0.008 toward positive, emission rate drops below 1.0 (abstention appears), and saved $\geq$ broke.

4. The Reward Problem, Explained Simply

Think of grading a student's "double-check" work:

- **PMI (current) $\approx$ "grade the double-check on how neat and natural it looks, as long as the final answer is right."** A student who writes a tidy "let me re-verify..." on every problem — even ones they already nailed — scores well. Worse, re-doing an already-correct problem occasionally introduces a fresh mistake. Technically PMI is the likelihood-delta: how much the meta text raises the probability of its own continuation versus a placebo ("Let me continue."), correctness-gated. It measures **fluency-correlated-with-correctness**, not whether the check did anything.
- **Counterfactual (proposed) $\approx$ "grade the double-check by comparing the student's score with it versus without it."** A double-check that

changes nothing earns nothing (so stop doing it on easy problems). A double-check that catches a real error earns a lot. A "double-check" that corrupts a right answer is penalized. This is exactly the behavior we want.

5. The Fix Being Implemented (design)

Backbone — group-branch counterfactual `R_meta` : split each 8-rollout group into 4 with-meta + 4 without-meta (meta-open token banned). Grade both arms against gold. The group-relative correctness advantage now embeds `correct_with - correct_without` at no extra decode cost. (Decisions locked: **4/4 split**, route onto the **answer region**, **fully replace** PMI.)

Shaping layered on (all from 2024-25 prior work, computed from the same branched siblings, no extra decode): - **SCoRe transition bonus** (2409.12917): `$\alpha \cdot (\text{correct\_with} - \text{correct\_without})$` with $\alpha > 1$ so right $\rightarrow$ wrong (breaking a correct answer) is the most-penalized event — the direct cure for "breaks > saves." - **AdaCoT over-trigger penalty + decision-token masking** (2505.11896): penalize emitting meta on a problem the without-meta arm already solved, and protect the verify/redirect/ none decision from collapsing into "always trigger." - **AdaptThink accuracy-floor** (2505.13417): forbid pushing meta-emission up on a group where with-meta accuracy is below without-meta — net-harmful meta is structurally blocked.

Safety: every new knob defaults OFF, so existing configs are byte-identical (no regression — test T1); a synthetic-group test (T2) proves the new heads actually reach the advantage computation (guarding against this project's recurring "computed-but-unused" inert-reward trap).

6. Current Status

- **Implementation:** an orchestrated multi-agent run (plan $\rightarrow$ TDD implement $\rightarrow$ adversarial verify $\rightarrow$ fix) is editing the core RL reward/advantage code (≈ 363 lines across `dcpo_region.py`, `verl_sdc.py`, `verl_sdc_utils.py`, the config) plus a new test file. Awaiting the adversarial-verify result before we commit.
- **Infrastructure notes (for reproducibility):** the original RL run hung mid-training at step 194 on a verl+vLLM deadlock (a healthy node, no code

fault); two resume attempts lost the Standard-tier idle-suspend lottery, so we pivoted to evaluating Δ at the saved gs190 checkpoint (90 steps past gs100) — which is what gave the decisive H1 result.

7. Limitations and Next Experiments

Limitations. - **Reward sparsity (the main risk of H3).** The counterfactual signal is informative only on mixed-correctness groups (where with/without arms disagree); easy (all-correct) and hard (all-wrong) groups give ~zero meta signal, and a 4-sample arm is high-variance. Mitigation: pair with **mixed-band training data** (filter to problems whose pass-rate is 0.125–0.875), so most groups are informative; fallbacks are a small dense PMI shaping term or a dense-causal PRIME value head (2502.01456). - H2 is a triangulated diagnosis, not a controlled ablation; H3's A/B is the ablation. - Δ is measured on a 594-problem held-out set; small strata (e.g. hard/verify, n=22) are noisy.

Next experiments. - **E1 (this fix):** run RL with group-branch counterfactual + SCoRe/AdaCoT shaping; judge by $\Delta > 0$, emission < 1.0 , saved $\geq$ broke, wellformed ≥ 0.40 . Monitor live with the autoresearch loop. - **E2 (sparsity):** the same reward on mixed-band data; compare Δ and meta-learning stability against E1. - **E3 (validation):** final 1030-problem @ 16k headline eval versus base-GRPO (0.770) and meta-GRPO (0.817); cross-model causal check (Spurious-Rewards safeguard, 2506.10947) to confirm $\Delta > 0$ is real and not a GRPO prior-amplification artifact.

References

- Math-Shepherd (2312.08935); Let's Verify Step by Step / PRM800K (2305.20050); Implicit PRM / Free Process Rewards (2412.01981); PRIME (2502.01456); GRPO-as-PRM / λ -GRPO (2509.21154 / 2510.00194); RLOO (2402.14740); SCoRe (2409.12917); AdaCoT (2505.11896); AdaptThink (2505.13417); Spurious Rewards in RLVR (2506.10947); RLCR (2507.16806).